

Assignment1_EE5178

October 23, 2022

1 Modern Computer Vision Assignment 1: Report

2 Importing Necessary Libraries

```
[1]: import torchvision
import numpy as np
import matplotlib.pyplot as plt
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import sklearn.metrics
import torch.nn as nn
import torch.nn.functional as F
import time
import sys
import os
import gc
import numpy as np
import PIL.Image
```

```
[2]: # Checking if cuda available
device=torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

3 Defining Hyperparameters and Necessary Global Variables

```
[3]: epochs = 10
learningRate = 0.01
inputs = 3*32*32
noOfNodesL1 = 500
noOfNodesL2 = 250
noOfNodesL3 = 100
```

```

outputs = 10
batchSize = 100
outputLabel = {0: 'airplane' , 1: 'automobile' , 2: 'bird' , 3: 'cat' , 4: 'deer' , 5: 'dog' , 6: 'frog' , 7: 'horse' , 8: 'ship' , 9: 'truck'}

```

4 Helpful Functions

```

[4]: def displayImages2(data , model ,device=torch.device("cuda" if torch.cuda.
      ↪is_available() else "cpu") ):
      """
      Displaying 'numberOfImages' Random Images
      numberOfImages - Even Number
      """
      ix=1
      fig, ax = plt.subplots(10, 5, figsize=(15,15))
      fig.tight_layout()
      model.eval()
      for filter in range(50):
          randInt = torch.randint(0 , 10000 , (1,))
          #print(randInt)
          #randInt = 1
          ax = plt.subplot(10, 5, ix)
          ax.set_xticks([])
          ax.set_yticks([])
          image , target = data[int(randInt)]
          temp = image.cpu()
          image = torch.reshape(image , (1 , 3 , 32 , 32))
          image = image.to(device)
          temp = torch.clamp(temp , 0 , 1)
          out = model(image)
          yPred = F.softmax(out , dim =1)
          predictedNumber = torch.argmax(yPred , dim = 1)
          ax.title.set_text('P | R: {}|{}'.format(outputLabel[int(predictedNumber)] ,
          ↪outputLabel[int(target)]))
          plt.imshow(temp.T)
          ix+=1

```

```

[15]: def loadDataSetPy(): # Loading DataSet for pyTorch
      """
      60k Train Data in batches of "batchSize"
      10k Test Data
      Also Does Appropriate Transformations Required
      """
      train_transform = transforms.Compose([
          transforms.ToTensor(),

```

```

transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),])

test_transform = transforms.Compose([
transforms.ToTensor(),
transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),])

train = torchvision.datasets.CIFAR10('./data' , train=True, transform=
→train_transform, download=True)

testData = torchvision.datasets.CIFAR10('./data' , train=False, transform=
→test_transform, download = True)
trainData, validData = torch.utils.data.random_split(train, [45000, 5000])

trainDataLoader = DataLoader(trainData , batch_size = batchSize , shuffle =
→True)
testDataLoader = DataLoader(trainData , batch_size = 10000 , shuffle = True)
validDataLoader = DataLoader(validData, batch_size = 5000 , shuffle = True ,
→num_workers = 2)

return trainDataLoader , testDataLoader , trainData, validDataLoader, testData

```

[104]: *# Defining the MLP*

```

class NeuralNet(nn.Module):
    def __init__(self):
        super(NeuralNet, self).__init__()
        self.relu = nn.ReLU()
        self.l1 = nn.Linear(inputs, noOfNodesL1)
        self.l2 = nn.Linear(noOfNodesL1, noOfNodesL2)
        self.l3 = nn.Linear(noOfNodesL2 , noOfNodesL3)
        self.l4 = nn.Linear(noOfNodesL3 , outputs)
    def forward(self, x):
        out = self.relu(self.l1(x))
        out = self.relu(self.l2(out))
        out = self.relu(self.l3(out))
        out = self.l4(out)
        # no activation and no softmax at the end
        return out

```

```

[6]: def creatingOutputVector(trainLabels):
    """
    trainLabels - Labels of Input Train Data of a Particular Batch. (Size:
    →batchSize X 1)
    yReal - True Output. Size: (batchSize X 10)
    """
    batchSize = len(trainLabels)

```

```

yReal = np.zeros([batchSize , 10]) # BatchSize X 10 Matrix
for j in range(batchSize):
    yReal[j][trainLabels[j]] = 1
return yReal

```

```

[106]: def predictPy(testDataLoader , model):
        for data in testDataLoader:
            img , label = data
            X = torch.reshape(img , (len(img) , 32*32*3))
            model.eval()
            yPred = F.softmax(model(X) , dim =1)
            predictedNumber = torch.argmax(yPred , dim = 1)
            noOfRightPrediction = torch.sum(predictedNumber == label)
            totalPrediction = len(img)
            accuracy = noOfRightPrediction/totalPrediction
            confusionMatrix = sklearn.metrics.confusion_matrix( label , predictedNumber)

        return accuracy , confusionMatrix

```

```

[5]: def predictPyCNN(testDataLoader , model , device=torch.device("cuda" if torch.
    ↪cuda.is_available() else "cpu")):
        if device == 'cuda':
            torch.cuda.empty_cache()
            gc.collect()
        for data in testDataLoader:
            img , label = data
            X = torch.reshape(img , (len(img) , 3 , 32 , 32))
            X = X.to(device)
            label = label.to(device)
            model.eval()
            yPred = F.softmax(model(X) , dim =1)
            predictedNumber = torch.argmax(yPred , dim = 1)
            noOfRightPrediction = torch.sum(predictedNumber == label)
            totalPrediction = len(img)
            accuracy = noOfRightPrediction/totalPrediction

            confusionMatrix = sklearn.metrics.confusion_matrix( label.cpu() ,
    ↪predictedNumber.cpu())

        return accuracy , confusionMatrix

```

```
[13]: def displayImages(data , model):
    """
    Displaying 32 Random Images
    """
    ix=1
    fig, ax = plt.subplots(8, 4, figsize=(10,10))
    fig.tight_layout()
    model.eval()
    for filter in range(32):
        randInt = torch.randint(0 , 10000 , (1,))
        ax = plt.subplot(8, 4, ix)
        ax.set_xticks([])
        ax.set_yticks([])
        image , target = data[int(randInt)]
        temp = image
        temp = torch.clamp(temp , 0 , 1)
        image = torch.reshape(image , (1 , 32*32*3))
        out = model(image)
        yPred = F.softmax(out , dim =1)
        predictedNumber = torch.argmax(yPred , dim = 1)
        ax.title.set_text('P | R: {}|{}'.format(outputLabel[int(predictedNumber)] ,
        ↪outputLabel[int(target)]))
        plt.imshow(temp.T)
        ix+=1
```

5 Training The Model

```
[16]: """
    Loading Datasets
    """
    trainDataLoader , testDataLoader , trainData, validDataLoader, testData =
    ↪loadDataSetPy()
```

Files already downloaded and verified

Files already downloaded and verified

```
[9]: """
    Intilizing the model
    """
    model = NeuralNet()
    optim = torch.optim.SGD(model.parameters(), lr=learningRate , momentum = 0.9)
    criterion = nn.CrossEntropyLoss()

    """
    Training
```

```

"""
start_time = time.time()
n_total_steps = batchSize
totalCost = []
model.train()
for epoch in range(epochs):
    print('Epoch {}'.format(epoch))
    for i, (images, labels) in enumerate(trainDataLoader):
        # FLATTENING
        images = torch.reshape(images, (batchSize, 3072)) # Size : [batchSize X 3072]
        # FORWARD PASS
        yPred = model(images)
        loss = criterion(yPred, labels)

        # OPTIMISATION
        optim.zero_grad()
        loss.backward()
        optim.step()
        cost = loss.item()
        if i%100 == 0:
            print('Cost After Mini Batch {}: {}'.format(i, cost))
            totalCost.append(loss.item())
    model.eval()
    accuracy, confusionMatrix = predictPy(validDataLoader, model)
    print('After Epoch {}, Validation accuracy = {}'.format(epoch, accuracy*100))
    plt.plot(totalCost)
    plt.title('After Epoch {}'.format(epoch))
    plt.xlabel('Mini Batches')
    plt.ylabel('Cost')
    plt.show()
print("Execution Time --- %s seconds ---" % (time.time() - start_time))

```

Epoch 0

Cost After Mini Batch 0: 2.311132

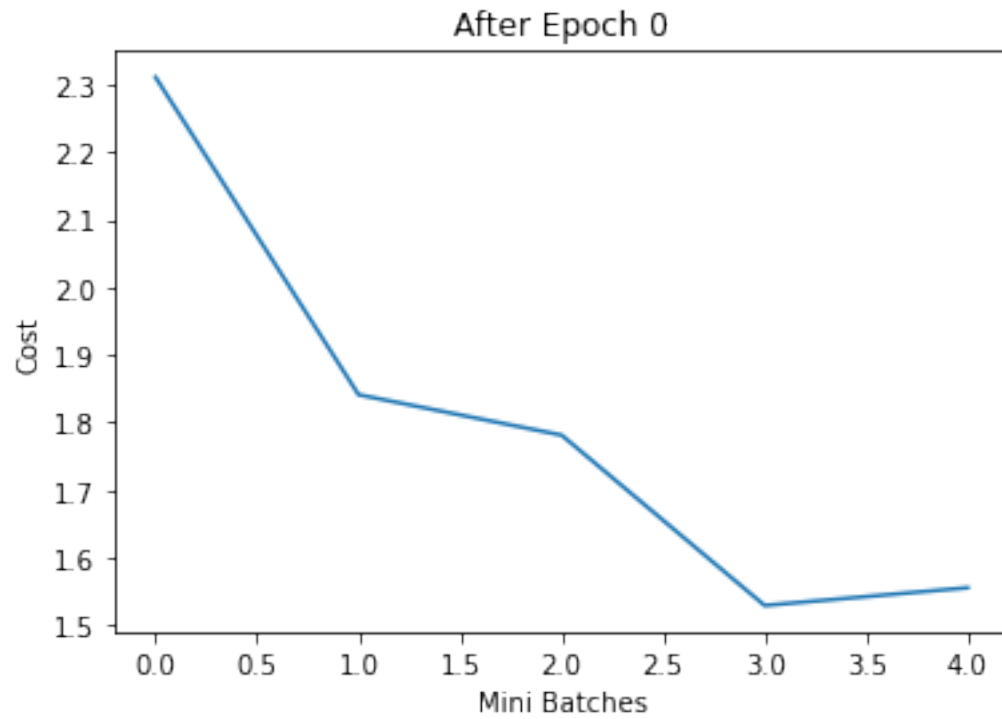
Cost After Mini Batch 100: 1.840939

Cost After Mini Batch 200: 1.781129

Cost After Mini Batch 300: 1.528851

Cost After Mini Batch 400: 1.555144

After Epoch 0, Validation accuracy = 45.41999816894531



Epoch 1

Cost After Mini Batch 0: 1.539053


Cost After Mini Batch 100: 1.522800

Cost After Mini Batch 200: 1.427155

Cost After Mini Batch 300: 1.646216

Cost After Mini Batch 400: 1.561159

After Epoch 1, Validation accuracy = 49.119998931884766



Epoch 2

Cost After Mini Batch 0: 1.298357


Cost After Mini Batch 100: 1.293602

Cost After Mini Batch 200: 1.274947

Cost After Mini Batch 300: 1.341696

Cost After Mini Batch 400: 1.179518

After Epoch 2, Validation accuracy = 51.279998779296875



Epoch 3

Cost After Mini Batch 0: 1.132018

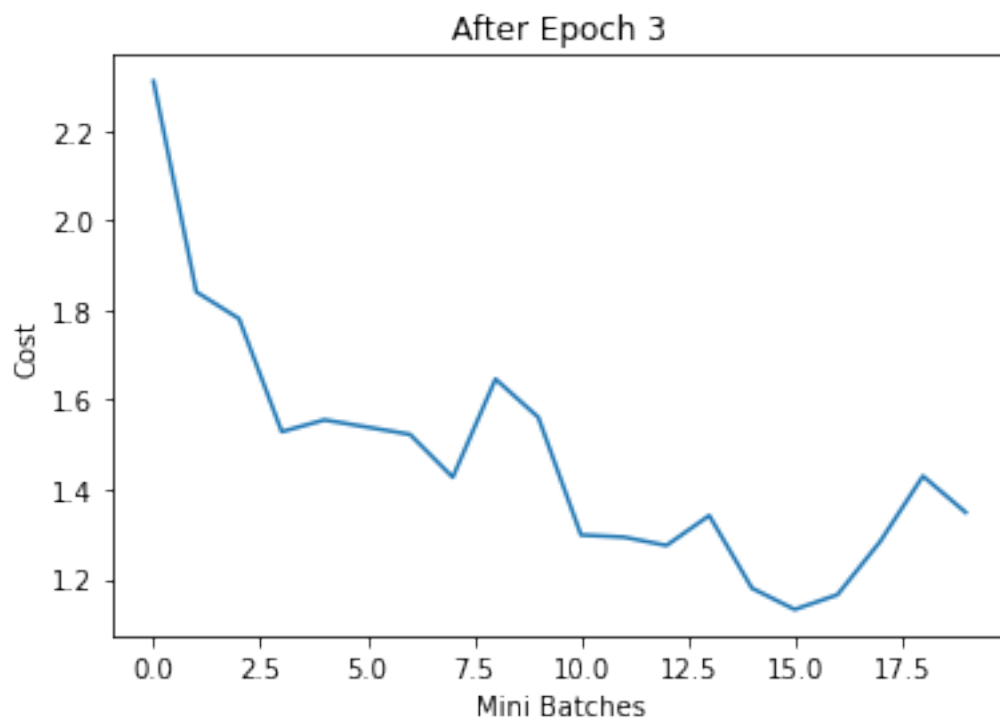
Cost After Mini Batch 100: 1.164940

Cost After Mini Batch 200: 1.283786

Cost After Mini Batch 300: 1.429813

Cost After Mini Batch 400: 1.348689

After Epoch 3, Validation accuracy = 51.12000274658203



Epoch 4

Cost After Mini Batch 0: 1.331268


Cost After Mini Batch 100: 1.043499

Cost After Mini Batch 200: 1.116550

Cost After Mini Batch 300: 0.997852

Cost After Mini Batch 400: 1.110326

After Epoch 4, Validation accuracy = 52.439998626708984



Epoch 5

Cost After Mini Batch 0: 1.114824


Cost After Mini Batch 100: 1.015516

Cost After Mini Batch 200: 0.964952

Cost After Mini Batch 300: 1.227359

Cost After Mini Batch 400: 1.299149

After Epoch 5, Validation accuracy = 52.70000076293945



Epoch 6

Cost After Mini Batch 0: 1.053469


Cost After Mini Batch 100: 1.139132

Cost After Mini Batch 200: 0.929680

Cost After Mini Batch 300: 0.970843

Cost After Mini Batch 400: 1.123901

After Epoch 6, Validation accuracy = 53.70000076293945



Epoch 7

Cost After Mini Batch 0: 0.813730

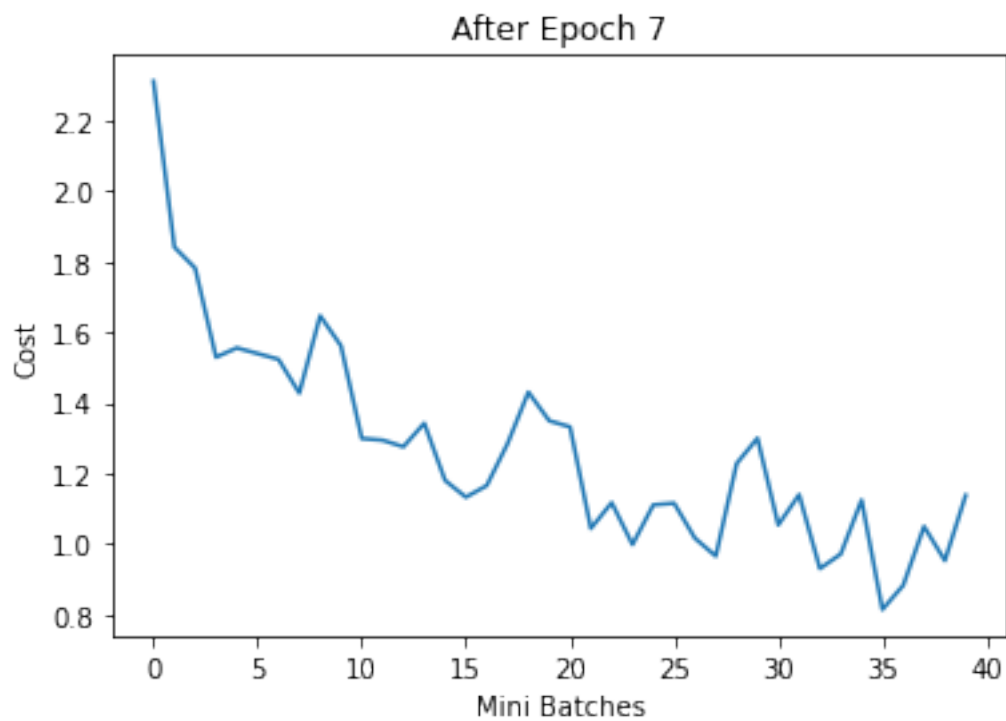
Cost After Mini Batch 100: 0.882246

Cost After Mini Batch 200: 1.048999

Cost After Mini Batch 300: 0.952442

Cost After Mini Batch 400: 1.137312

After Epoch 7, Validation accuracy = 52.78000259399414



Epoch 8

Cost After Mini Batch 0: 0.865818

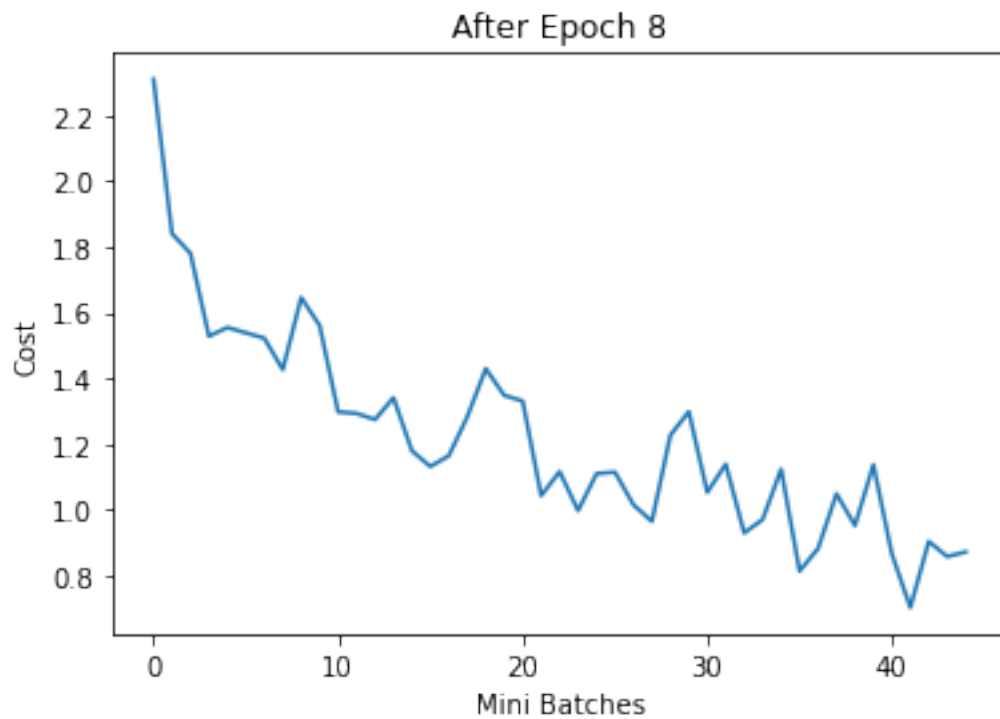
Cost After Mini Batch 100: 0.702905

Cost After Mini Batch 200: 0.904021

Cost After Mini Batch 300: 0.857858

Cost After Mini Batch 400: 0.872158

After Epoch 8, Validation accuracy = 52.619998931884766



Epoch 9

Cost After Mini Batch 0: 0.748782


Cost After Mini Batch 100: 0.712571

Cost After Mini Batch 200: 0.846918

Cost After Mini Batch 300: 0.804748

Cost After Mini Batch 400: 1.026457

After Epoch 9, Validation accuracy = 53.500003814697266



Execution Time --- 171.6047477722168 seconds ---

6 Confusion Matrix

```
[10]: accuracy , confusionMatrix = predictPy(testDataLoader , model)
print('Accuracy on Test Data = {}'.format(accuracy*100))
print('Confusion Matrix on Test Data')
print(confusionMatrix)
```

Accuracy on Test Data = 73.6199951171875%

Confusion Matrix on Test Data

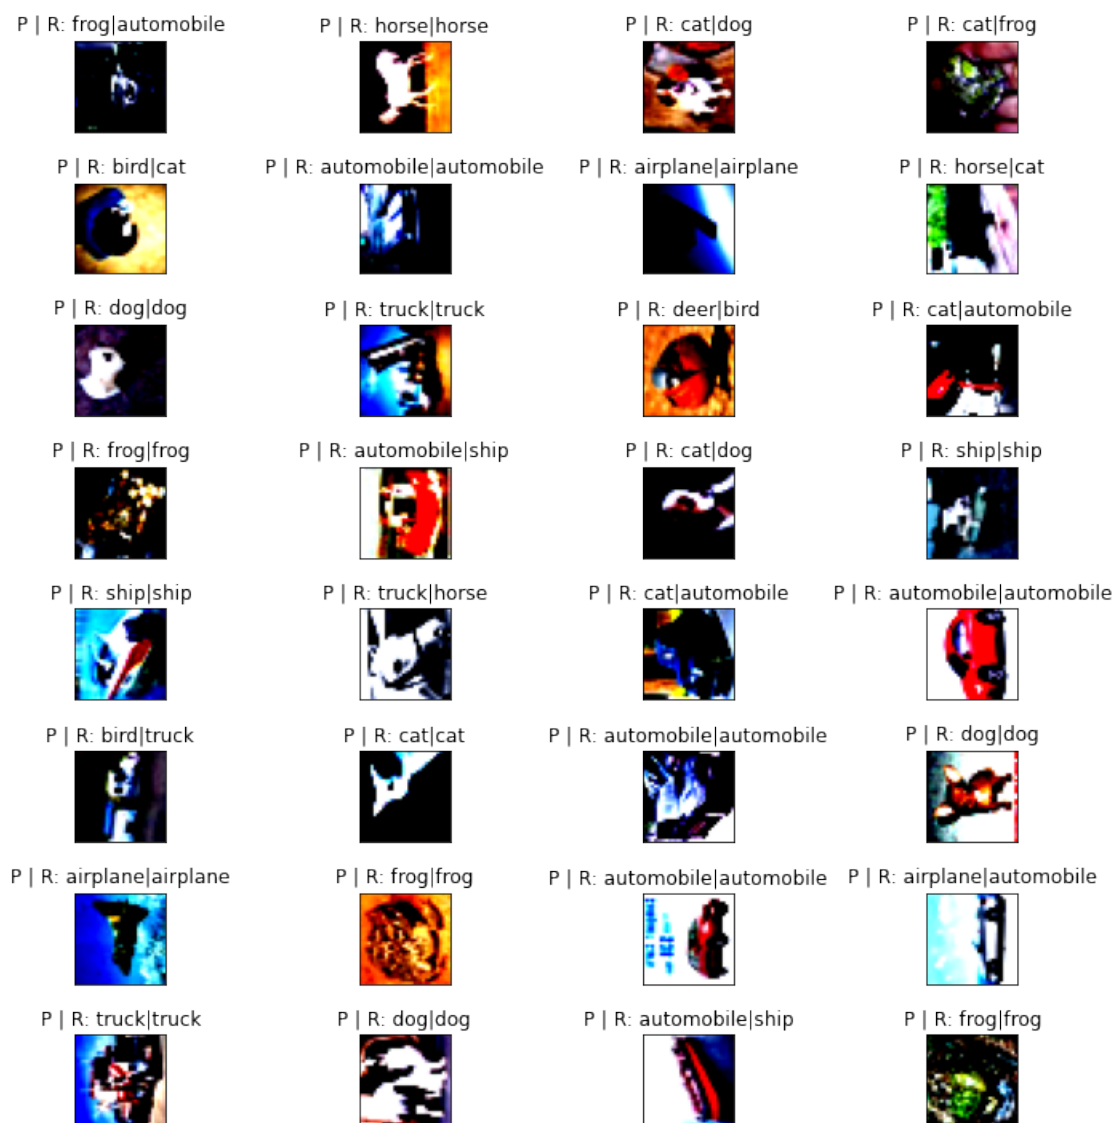
```
[[349 10 27 12 24 4 4 3 43 22]
 [ 3 434 5 7 6 3 6 2 24 45]
 [ 21 1 263 50 67 23 27 9 5 2]
 [ 7 2 28 306 26 85 23 11 6 6]
 [ 11 5 28 21 383 4 37 11 6 3]
 [ 5 1 15 82 27 347 14 8 1 3]
 [ 2 3 13 41 28 8 394 3 0 3]
 [ 5 2 6 25 57 22 3 369 4 8]
 [ 21 27 4 11 8 1 2 0 423 12]
 [ 5 29 3 11 6 4 2 3 6 413]]
```


7 Plotting Randomly Selected Images

- Without Batch Normalization
- P : Predicted Class, R : Real Class

```
[12]: displayImages(testData , model)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:25: UserWarning:
The use of `x.T` on tensors of dimension other than 2 to reverse their shape is
deprecated and it will throw an error in a future release. Consider `x.mT` to
transpose batches of matrices or `x.permute(*torch.arange(x.ndim - 1, -1, -1))`
to reverse the dimensions of a tensor. (Triggered internally at
../aten/src/ATen/native/TensorShape.cpp:2981.)
```



8 With Batch Normalization

```
[ ]: class NeuralNetWithBN(nn.Module):
    def __init__(self):
        super(NeuralNetWithBN, self).__init__()
        self.relu = nn.ReLU()
        self.l1 = nn.Linear(inputs, noOfNodesL1)
        self.b1 = nn.BatchNorm1d(noOfNodesL1)
        self.l2 = nn.Linear(noOfNodesL1, noOfNodesL2)
        self.b2 = nn.BatchNorm1d(noOfNodesL2)
        self.l3 = nn.Linear(noOfNodesL2 , noOfNodesL3)
        self.b3 = nn.BatchNorm1d(noOfNodesL3)
        self.l4 = nn.Linear(noOfNodesL3 , outputs)
    def forward(self, x):
        out = self.relu(self.b1(self.l1(x)))
        out = self.relu(self.b2(self.l2(out)))
        out = self.relu(self.b3(self.l3(out)))
        out = self.l4(out)
        # no activation and no softmax at the end
        return out

[ ]: """
Intilizing the model
"""
modelWithBN = NeuralNetWithBN()
optim = torch.optim.SGD(modelWithBN.parameters(), lr=learningRate , momentum = 0.9)
criterion = nn.CrossEntropyLoss()

"""
Training
"""
modelWithBN.train()
start_time = time.time()
n_total_steps = batchSize
totalCost = []
for epoch in range(epochs):
    print('Epoch {}'.format(epoch))
    for i , (images,labels) in enumerate(trainDataLoader):
        # FLATTENING
        images = torch.reshape(images , (batchSize , 3072)) # Size : [batchSize X 3072]
        # FORWARD PASS
        yPred = modelWithBN(images)
        loss = criterion(yPred , labels)

    # OPTIMISATION
```

```

optim.zero_grad()
loss.backward()
optim.step()
cost = loss.item()
if i%100 == 0:
    print('Cost After Mini Batch %i: %f' %(i , cost))
    totalCost.append(loss.item())
modelWithBN.eval()
accuracy, confusionMatrix = predictPy(validDataLoader , modelWithBN)
print('After Epoch {}, Validation accuracy = {}'.format(epoch , accuracy*100))
plt.plot(totalCost)
plt.title('After Epoch {}'.format(epoch))
plt.xlabel('Mini Batches')
plt.ylabel('Cost')
plt.show()
print("Execution Time --- %s seconds ---" % (time.time() - start_time))

```

Epoch 0

Cost After Mini Batch 0: 2.383224


Cost After Mini Batch 100: 1.700450

Cost After Mini Batch 200: 1.453422

Cost After Mini Batch 300: 1.522559

Cost After Mini Batch 400: 1.391195

After Epoch 0, Validation accuracy = 48.5



Epoch 1

Cost After Mini Batch 0: 1.582003


Cost After Mini Batch 100: 1.543284

Cost After Mini Batch 200: 1.427740

Cost After Mini Batch 300: 1.199081

Cost After Mini Batch 400: 1.192501

After Epoch 1, Validation accuracy = 48.97999954223633



Epoch 2

Cost After Mini Batch 0: 1.375597


Cost After Mini Batch 100: 1.322512

Cost After Mini Batch 200: 1.235305

Cost After Mini Batch 300: 1.281783

Cost After Mini Batch 400: 1.109986

After Epoch 2, Validation accuracy = 49.63999938964844



Epoch 3

Cost After Mini Batch 0: 1.221174

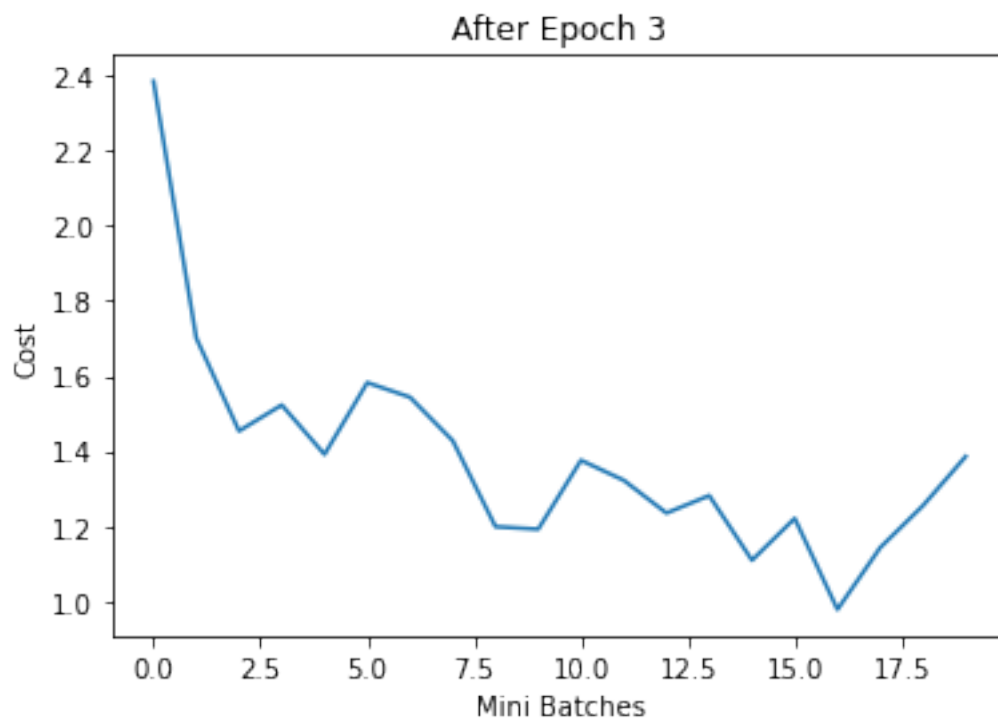
Cost After Mini Batch 100: 0.979037

Cost After Mini Batch 200: 1.143501

Cost After Mini Batch 300: 1.254964

Cost After Mini Batch 400: 1.385743

After Epoch 3, Validation accuracy = 52.47999954223633



Epoch 4

Cost After Mini Batch 0: 1.100049

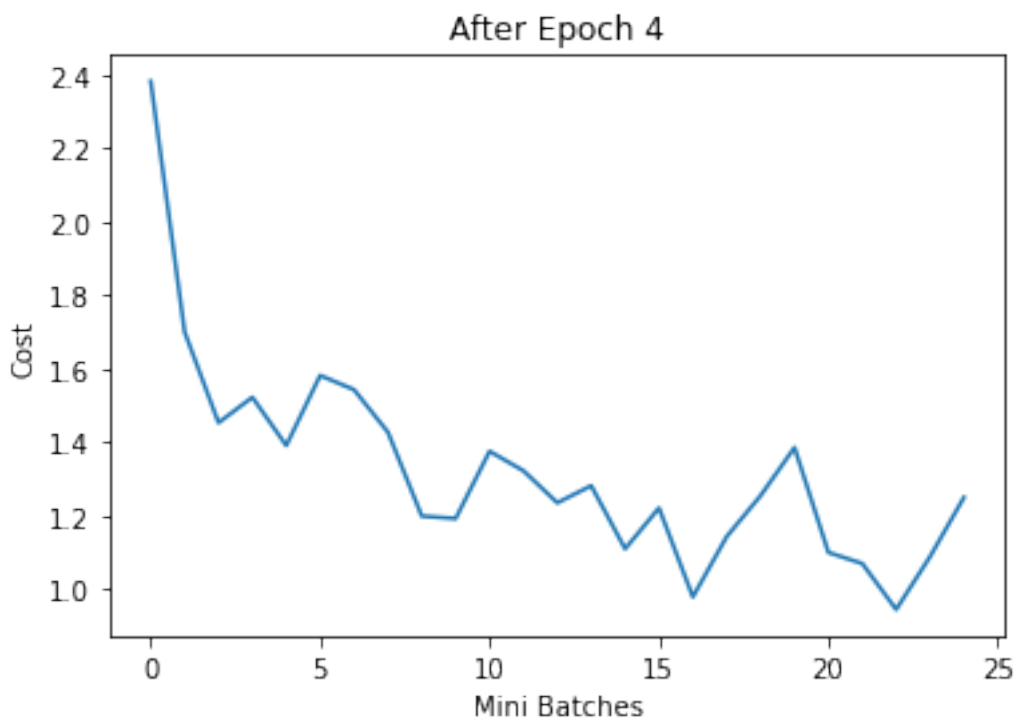
Cost After Mini Batch 100: 1.069984

Cost After Mini Batch 200: 0.944833

Cost After Mini Batch 300: 1.088776

Cost After Mini Batch 400: 1.249971

After Epoch 4, Validation accuracy = 52.439998626708984



Epoch 5

Cost After Mini Batch 0: 1.051230


Cost After Mini Batch 100: 1.046620

Cost After Mini Batch 200: 1.222595

Cost After Mini Batch 300: 1.061490

Cost After Mini Batch 400: 1.140897

After Epoch 5, Validation accuracy = 55.159996032714844



Epoch 6

Cost After Mini Batch 0: 0.843136


Cost After Mini Batch 100: 1.013773

Cost After Mini Batch 200: 0.947004

Cost After Mini Batch 300: 0.942487

Cost After Mini Batch 400: 0.951548

After Epoch 6, Validation accuracy = 54.07999801635742



Epoch 7

Cost After Mini Batch 0: 0.792948


Cost After Mini Batch 100: 0.656006

Cost After Mini Batch 200: 1.014510

Cost After Mini Batch 300: 1.006602

Cost After Mini Batch 400: 0.933423

After Epoch 7, Validation accuracy = 53.820003509521484



Epoch 8

Cost After Mini Batch 0: 0.858434

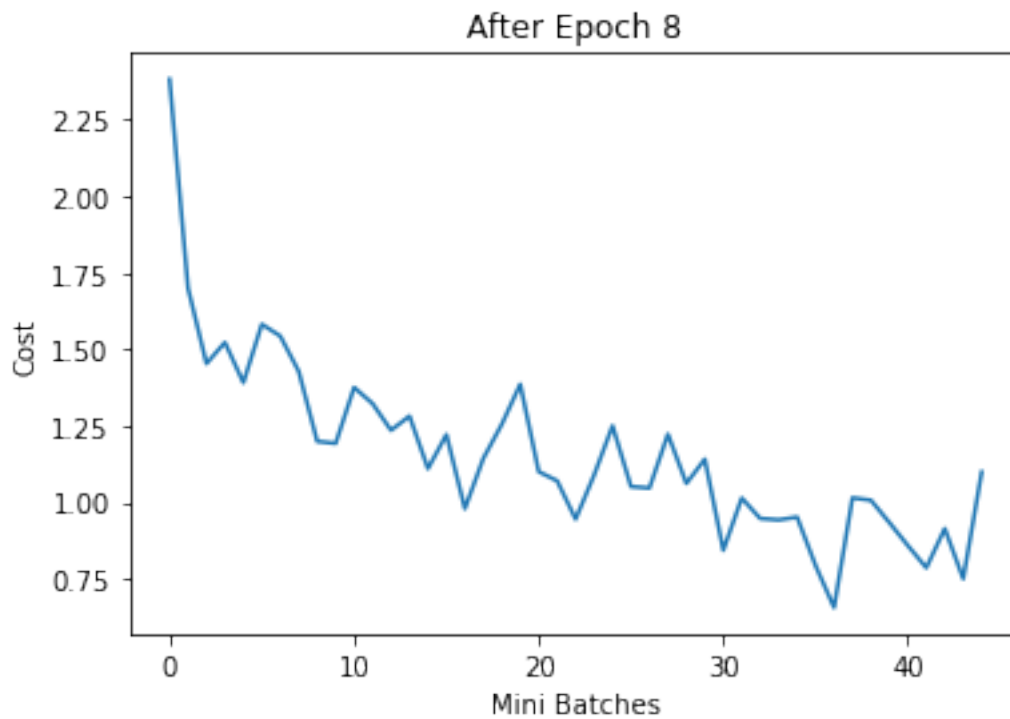
Cost After Mini Batch 100: 0.785785

Cost After Mini Batch 200: 0.913952

Cost After Mini Batch 300: 0.749897

Cost After Mini Batch 400: 1.099093

After Epoch 8, Validation accuracy = 53.15999984741211



Epoch 9

Cost After Mini Batch 0: 0.894441


Cost After Mini Batch 100: 0.662178

Cost After Mini Batch 200: 0.641659

Cost After Mini Batch 300: 0.810018

Cost After Mini Batch 400: 0.655432

After Epoch 9, Validation accuracy = 53.96000289916992



Execution Time --- 225.135009765625 seconds ---

9 Confusion Matrix

```
[ ]: accuracy , confusionMatrix = predictPy(testDataLoader , modelWithBN)
print('Accuracy on Test Data = {}'.format(accuracy*100))
print('Confusion Matrix on Test Data')
print(confusionMatrix)
```

Accuracy on Test Data = 76.5%

Confusion Matrix on Test Data

```
[[437  11  13   2   5   3   5   8  21  10]
 [ 13 401   5   1   2   2   6   0   6  43]
 [ 29   5 302  34  34  16  33  12   3   4]
 [ 13   9  27 337  14  77  36  17   2   7]
 [ 17   5  41  17 350  13  22  23   1   2]
 [  5   3  26  81   9 340   8  24   1   1]
 [  4   1   8  21  13   5 443   1   0   1]
 [ 11   2  10  12  12  26   5 420   2   3]
 [ 72  24  10   5   7   7   3   2 368  12]
 [ 14  22   5   5   4   8   1   5   5 427]]
```

```
[ ]: displayImages(testData , modelWithBN)
```



10 Report

1. Prediction Accuracy on Test Data (Without Batch Normalization) - 73.6%

Prediction Accuracy on Test Data (With Batch Normalization) - 76.5 %

2. **Use batch-normalization. Does it improve the test accuracy? Does it affect training time?**

Yes as seen, from above, Batch Normalization has improved the accuracy by a small margin.

Maybe this can also be associated with the initialization of weights.

Execution Time

With Batch Normalization - 225.14 secs

Without BN - 171.60 secs

We can see a considerable increase in Execution Time

11 CNN

Loading Datasets and Performing Necessary Transformations

```
[6]: train_transform = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

train_dset = torchvision.datasets.CIFAR10(root="data/", train=True,
    ↪transform=train_transform, download=True)
test_dset = torchvision.datasets.CIFAR10(root="data/", train=False,
    ↪transform=test_transform, download=True)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
data/cifar-10-python.tar.gz

0%| | 0/170498071 [00:00<?, ?it/s]

Extracting data/cifar-10-python.tar.gz to data/
Files already downloaded and verified

```
[7]: train_loader = DataLoader(train_dset, batch_size=100, shuffle=True,
    ↪num_workers=2)
test_loader = DataLoader(test_dset, batch_size=100, shuffle=False,
    ↪num_workers=2)
```

Defining the Architecture

```
[8]: class VGG(nn.Module):
    CONFIGS = {
```

```

        "vgg11": [64, "pool", 128, "pool", 256, 256, "pool", 512, 512, "pool",
→512, 512, "pool"],
    }

    def __init__(self, cfg):
        super(VGG, self).__init__()
        in_dim = 3
        layers = []
        for layer in self.CONFIGS[cfg]:
            if layer == "pool":
                maxpool = nn.MaxPool2d(kernel_size = 2 , stride = 2)
                layers.append(maxpool)
            else:
                block = nn.Sequential(nn.Conv2d(in_dim , layer , kernel_size = 3 ,
→padding = 1),
                                     nn.BatchNorm2d(layer),
                                     nn.ReLU())
                layers.append(block)
                in_dim = layer
        avgpool = nn.AvgPool2d(kernel_size=1)
        layers.append(avgpool)
        self.layers = nn.Sequential(*layers)
        self.fc1 = nn.Linear(512, 100)
        self.fc2 = nn.Linear(100,10)

    def forward(self, x):
        out = self.layers(x)
        out = torch.flatten(out , start_dim = 1)
        out = self.fc1(out)
        out = self.fc2(out)
        return out

```

```

[9]: def pbar(p=0, msg="", bar_len=20):
    sys.stdout.write("\033[K")
    sys.stdout.write("\x1b[2K" + "\r")
    block = int(round(bar_len * p))
    text = "Progress: [{}] {}% {}".format(
        "\x1b[32m" + "=" * (block - 1) + ">" + "\033[0m" + "-" * (bar_len -
→block),
        round(p * 100, 2),
        msg,
    )
    print(text, end="\r")
    if p == 1:
        print()

class AvgMeter:

```

```

def __init__(self):
    self.reset()

def reset(self):
    self.metrics = {}

def add(self, batch_metrics):
    if self.metrics == {}:
        for key, value in batch_metrics.items():
            self.metrics[key] = [value]
    else:
        for key, value in batch_metrics.items():
            self.metrics[key].append(value)

def get(self):
    return {key: np.mean(value) for key, value in self.metrics.items()}

def msg(self):
    avg_metrics = {key: np.mean(value) for key, value in self.metrics.
→items()}
    return "".join(["[{}] {:.5f} ".format(key, value) for key, value in
→avg_metrics.items()])

```

```

[10]: def train(model, optim, lr_sched=None, epochs=200, device=torch.device("cuda"
→if torch.cuda.is_available() else "cpu"), criterion=None, metric_meter=None,
→out_dir="out/"):
    if device == 'cuda':
        torch.cuda.empty_cache()
        gc.collect()
    model.to(device)
    best_acc = 0
    trainError = []
    validError = []
    predAcc = []
    for epoch in range(epochs):
        model.train()
        metric_meter.reset()
        for indx, (img, target) in enumerate(train_loader):
            img = img.to(device)
            target = target.to(device)

            out = model(img)
            loss = criterion(out, target)
            optim.zero_grad()
            loss.backward()
            optim.step()
            cost = loss.item()

```



```

    if indx % 5 == 0:
        trainError.append(cost)

    metric_meter.add({"train loss": cost})
    pbar(indx / len(train_loader), msg=metric_meter.msg())
    print('Please Work')
    pbar(1, msg=metric_meter.msg())

model.eval()
metric_meter.reset()
for indx, (img, target) in enumerate(test_loader):
    img = img.to(device)
    target = target.to(device)

    out= model(img)
    loss = criterion(out, target)
    cost = loss.item()
    if indx % 5 == 0:
        validError.append(cost)
        yPred = F.softmax(out , dim =1)
        predictedNumber = torch.argmax(yPred , dim = 1)
        noOfRightPrediction = torch.sum(predictedNumber == target)
        totalPrediction = yPred.shape[0]
        acc = noOfRightPrediction/totalPrediction
        acc = acc.cpu().detach().numpy()
        if indx % 5 == 0:
            predAcc.append(acc)

    metric_meter.add({"test loss": loss.item(), "test acc": acc})
    pbar(indx / len(test_loader), msg=metric_meter.msg())
    pbar(1, msg=metric_meter.msg())

test_metrics = metric_meter.get()
if test_metrics["test acc"] > best_acc:
    print(
        "\x1b[33m"
        + f"test acc improved from {round(best_acc, 5)} to_
→{round(test_metrics['test acc'], 5)}"
        + "\033[0m"
    )
    best_acc = test_metrics['test acc']
    torch.save(model.state_dict(), os.path.join(out_dir, "best.ckpt"))
    lr_sched.step()
plt.plot(trainError)
plt.xlabel('Epochs')
plt.ylabel('Training Error')
plt.show()

```

```

plt.plot(validError)
plt.xlabel('Epochs')
plt.ylabel('Validation Error')
plt.show()
plt.plot(predAcc)
plt.xlabel('Epochs')
plt.ylabel('Prediction Accuracy')
plt.show()

```

```

[11]: def run_experiment(epochs=200):
    model_name = "VGG"
    model_cfg = "vgg11"
    model = VGG('vgg11')
    optim = torch.optim.SGD(model.parameters(), lr=1e-1, momentum=0.9,
    ↪weight_decay=5e-4)
    lr_sched = torch.optim.lr_scheduler.CosineAnnealingLR(optim, T_max=epochs)
    criterion = nn.CrossEntropyLoss()
    metric_meter = AvgMeter()
    out_dir = f"{model_name}_{model_cfg}"
    os.makedirs(out_dir, exist_ok=True)
    train(model, optim, lr_sched, epochs=epochs, criterion=criterion,
    ↪metric_meter=metric_meter, out_dir=out_dir)
    return model

```

12 Training the CNN

Displayed the Train Cost, Validation Cost and Validation Accuracy

```

[12]: modelVGG = run_experiment(epochs=20)

```

```

Please Work
Progress: [=====>] 100% [train loss] 4.37105
Progress: [=====>] 100% [test loss] 2.33205 [test acc]
0.10000
test acc improved from 0 to 0.10000000149011612
Please Work
Progress: [=====>] 100% [train loss] 2.33519
Progress: [=====>] 100% [test loss] 2.31161 [test acc]
0.10000
Please Work
Progress: [=====>] 100% [train loss] 2.30926
Progress: [=====>] 100% [test loss] 2.31434 [test acc]
0.10090
test acc improved from 0.10000000149011612 to 0.10090000182390213
Please Work
Progress: [=====>] 100% [train loss] 2.29586

```

```



Progress: [=====>] 100% [test loss] 2.28165 [test acc]
0.12980
test acc improved from 0.10090000182390213 to 0.1298000067472458
Please Work
Progress: [=====>] 100% [train loss] 2.03508
Progress: [=====>] 100% [test loss] 1.81641 [test acc]
0.28450
test acc improved from 0.1298000067472458 to 0.28450000286102295
Please Work
Progress: [=====>] 100% [train loss] 1.77731
Progress: [=====>] 100% [test loss] 1.67264 [test acc]
0.33160
test acc improved from 0.28450000286102295 to 0.33160001039505005
Please Work
Progress: [=====>] 100% [train loss] 1.60010
Progress: [=====>] 100% [test loss] 1.56384 [test acc]
0.40560
test acc improved from 0.33160001039505005 to 0.40560001134872437
Please Work
Progress: [=====>] 100% [train loss] 1.40590
Progress: [=====>] 100% [test loss] 1.28519 [test acc]
0.52850
test acc improved from 0.40560001134872437 to 0.5285000205039978
Please Work
Progress: [=====>] 100% [train loss] 1.23223
Progress: [=====>] 100% [test loss] 1.49581 [test acc]
0.48380
Please Work
Progress: [=====>] 100% [train loss] 1.09219
Progress: [=====>] 100% [test loss] 1.06945 [test acc]
0.62690
test acc improved from 0.5285000205039978 to 0.6269000172615051
Please Work
Progress: [=====>] 100% [train loss] 0.98541
Progress: [=====>] 100% [test loss] 1.15533 [test acc]
0.62230
Please Work
Progress: [=====>] 100% [train loss] 0.90046
Progress: [=====>] 100% [test loss] 0.91588 [test acc]
0.69180
test acc improved from 0.6269000172615051 to 0.6917999982833862
Please Work
Progress: [=====>] 100% [train loss] 0.82347
Progress: [=====>] 100% [test loss] 0.80108 [test acc]
0.72270
test acc improved from 0.6917999982833862 to 0.722699998092651
Please Work
Progress: [=====>] 100% [train loss] 0.74968


```

```

Progress: [=====>] 100% [test loss] 0.81656 [test acc]
0.71700
Please Work
Progress: [=====>] 100% [train loss] 0.68185
Progress: [=====>] 100% [test loss] 0.69564 [test acc]
0.76420
test acc improved from 0.7226999998092651 to 0.76419997215271
Please Work
Progress: [=====>] 100% [train loss] 0.61270
Progress: [=====>] 100% [test loss] 0.63950 [test acc]
0.78480
test acc improved from 0.76419997215271 to 0.7847999930381775
Please Work
Progress: [=====>] 100% [train loss] 0.55754
Progress: [=====>] 100% [test loss] 0.60501 [test acc]
0.79420
test acc improved from 0.7847999930381775 to 0.7942000031471252
Please Work
Progress: [=====>] 100% [train loss] 0.50129
Progress: [=====>] 100% [test loss] 0.56405 [test acc]
0.80670
test acc improved from 0.7942000031471252 to 0.8066999912261963
Please Work
Progress: [=====>] 100% [train loss] 0.46348
Progress: [=====>] 100% [test loss] 0.52243 [test acc]
0.82490
test acc improved from 0.8066999912261963 to 0.8248999714851379
Please Work
Progress: [=====>] 100% [train loss] 0.43607
Progress: [=====>] 100% [test loss] 0.50315 [test acc]
0.83190
test acc improved from 0.8248999714851379 to 0.8319000005722046

```





13 Confusion Matrix

```
[17]: accuracy , confusionMatrix = predictPyCNN(testDataLoader , modelVGG ,
↪device=torch.device("cuda" if torch.cuda.is_available() else "cpu"))
print('Accuracy on Test Data = {}'.format(accuracy*100))
print('Confusion Matrix on Test Data')
print(confusionMatrix)
```


Accuracy on Test Data = 87.15999603271484%

Confusion Matrix on Test Data

```
[[444  1  7  1  3  1  0  2 28  5]
 [  3 469  0  0  0  1  1  0  5 17]
 [ 15  1 378 18 17 11 12  6  1  2]
 [  9  1 17 367 12 73 22  8  0  5]
 [  3  0 15 19 428  6  7 15  2  0]
 [  1  1  6 61 14 397  6 13  0  0]
 [  0  0 10 15 15  5 465  1  2  1]
 [  2  0  4 18  8 14  1 447  0  2]
 [ 10  2  1  6  1  0  0  1 480  5]
 [ 10 26  0  1  0  0  1  2  4 483]]
```

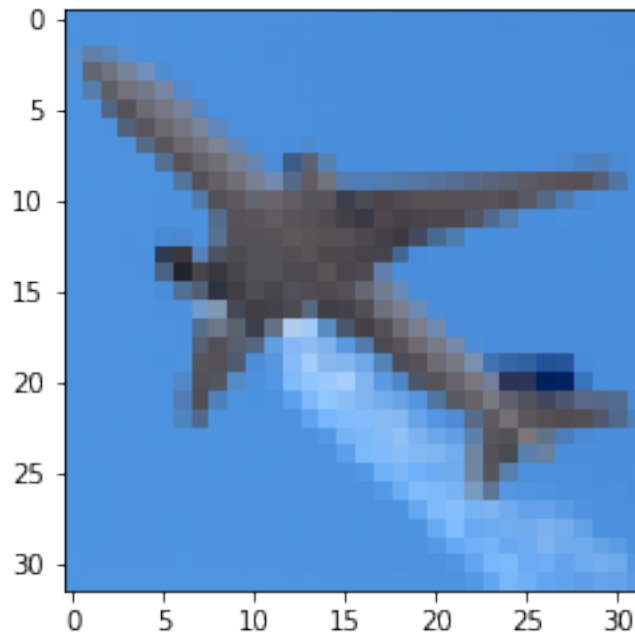
```
[34]: displayImages2(testData , modelVGG)
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:26: UserWarning:
The use of `x.T` on tensors of dimension other than 2 to reverse their shape is
deprecated and it will throw an error in a future release. Consider `x.mT` to
transpose batches of matrices or `x.permute(*torch.arange(x.ndim - 1, -1, -1))`
to reverse the dimensions of a tensor. (Triggered internally at
../aten/src/ATen/native/TensorShape.cpp:2981.)



14 Passing 5 Images downloaded from Internet

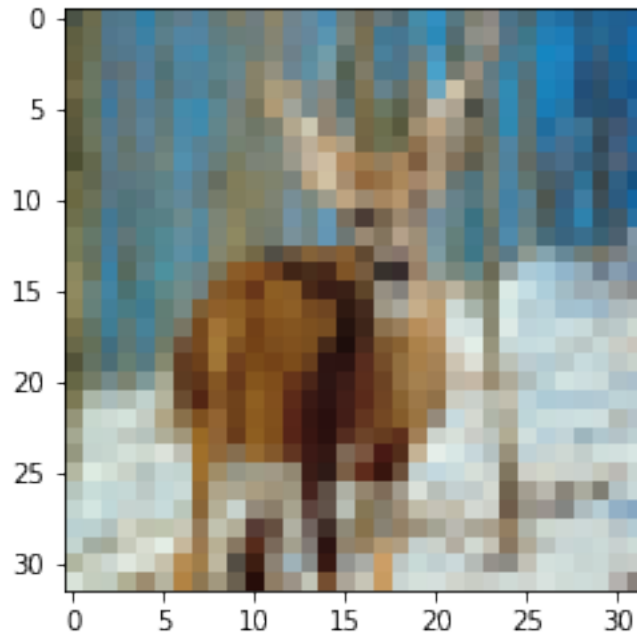
```
[19]: airplane = PIL.Image.open("/content/Airplane.png")
airplane = np.array(airplane.convert('RGB'))
plt.imshow(airplane)
plt.show()
```



```
[22]: image = test_transform(airplane)
temp = image.cpu()
image = torch.reshape(image , (1 , 3 , 32 , 32)).float()
image = image.to(device)
temp = torch.clamp(temp , 0 , 1)
out = modelVGG(image)
yPred = F.softmax(out , dim =1)
predictedNumber = torch.argmax(yPred , dim = 1)
print('Predicted Class -', outputLabel[int(predictedNumber)])
```

Predicted Class - airplane


```
[23]: deer = PIL.Image.open("/content/Deer.png")
deer = np.array(deer.convert('RGB'))
plt.imshow(deer)
plt.show()
```

```
[24]: image = test_transform(deer)
temp = image.cpu()
image = torch.reshape(image , (1 , 3 , 32 , 32)).float()
image = image.to(device)
temp = torch.clamp(temp , 0 , 1)
out = modelVGG(image)
yPred = F.softmax(out , dim =1)
predictedNumber = torch.argmax(yPred , dim = 1)
print('Predicted Class -', outputLabel[int(predictedNumber)])
```

Predicted Class - deer


```
[25]: frog = PIL.Image.open("/content/Frog.png")
frog = np.array(frog.convert('RGB'))
plt.imshow(frog)
plt.show()
```



```
[26]: image = test_transform(frog)
temp = image.cpu()
image = torch.reshape(image , (1 , 3 , 32 , 32)).float()
image = image.to(device)
temp = torch.clamp(temp , 0 , 1)
out = modelVGG(image)
yPred = F.softmax(out , dim =1)
predictedNumber = torch.argmax(yPred , dim = 1)
print('Predicted Class -', outputLabel[int(predictedNumber)])
```

Predicted Class - frog

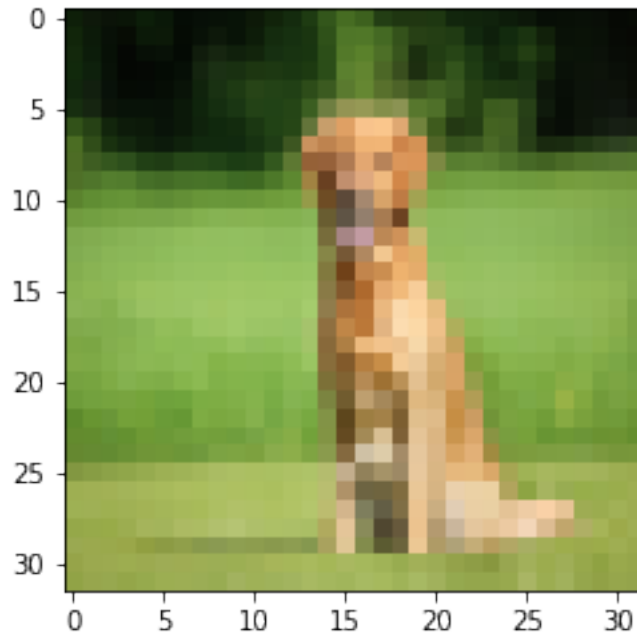
```
[27]: cat = PIL.Image.open("/content/cat.png")
cat = np.array(cat.convert('RGB'))
plt.imshow(cat)
plt.show()
```



```
[28]: image = test_transform(cat)
temp = image.cpu()
image = torch.reshape(image , (1 , 3 , 32 , 32)).float()
image = image.to(device)
temp = torch.clamp(temp , 0 , 1)
out = modelVGG(image)
yPred = F.softmax(out , dim =1)
predictedNumber = torch.argmax(yPred , dim = 1)
print('Predicted Class -', outputLabel[int(predictedNumber)])
```

Predicted Class - cat

```
[32]: dog = PIL.Image.open("/content/dog.png")
dog = np.array(dog.convert('RGB'))
plt.imshow(dog)
plt.show()
```



```
[33]: image = test_transform(dog)
temp = image.cpu()
image = torch.reshape(image , (1 , 3 , 32 , 32)).float()
image = image.to(device)
temp = torch.clamp(temp , 0 , 1)
out = modelVGG(image)
yPred = F.softmax(out , dim =1)
predictedNumber = torch.argmax(yPred , dim = 1)
print('Predicted Class -', outputLabel[int(predictedNumber)])
```

Predicted Class - deer

15 Observation

We can observe, almost all the random images were predicted correctly.